



Universidade do Minho
Escola de Engenharia
Departamento de Informática

Mestrado Integrado em Engenharia Informática

Computação Natural


2018/2019

Paulo Novais, Cesar Analide, Filipe Gonçalves

- Paulo Novais – pjon@di.uminho.pt
- Cesar Analide – cesar.analide@di.uminho.pt
- Filipe Gonçalves – fgoncalves@algoritmi.uminho.pt

- Departamento de Informática
Escola de Engenharia
Universidade do Minho
- Grupo ISLab – (Synthetic Intelligence Lab)
- Centro ALGORITMI
Universidade do Minho

Introduction to Reinforcement Learning



- The Multi-armed Bandit Problem (Exploration vs Exploitation Problem)



D1



D2



D3



D4



D5


- The Multi-armed Bandit Problem (Exploration vs Exploitation Problem)



■ Marketing - Ads Selection (Exploration vs Exploitation Problem)



D1




D2



D3



D4



D5

■ Marketing - Ads Selection (Exploration vs Exploitation Problem)

- We have d arms. For example, arms are ads that we display to users each time they connect to a web page
- Each time a user connects to this web page, that makes a round
- At each round n , we choose one ad to display to the user
- At each round n , ad i gives reward:
$$r_i(n) \in \{0, 1\}: r_i(n) = 1 \text{ if the user clicked on the ad } i, 0 \text{ if the user didn't}$$
- Goal: maximize the total reward we get over many rounds

■ Marketing - Ads Selection

- Index: Person
- Column: Ads
 - 1: Person clicked on Ad
 - 0: Person ignored Ad

Index	Ad 1	Ad 2	Ad 3	Ad 4	Ad 5	Ad 6	Ad 7	Ad 8	Ad 9	Ad 10
0	1	0	0	0	1	0	0	0	1	0
1	0	0	0	0	0	0	0	0	1	0
2	0	0	0	0	0	0	0	0	0	0
3	0	1	0	0	0	0	0	1	0	0
4	0	0	0	0	0	0	0	0	0	0
5	1	1	0	0	0	0	0	0	0	0
6	0	0	0	1	0	0	0	0	0	0
7	1	1	0	0	1	0	0	0	0	0
8	0	0	0	0	0	0	0	0	0	0
9	0	0	1	0	0	0	0	0	0	0
10	0	0	0	0	0	0	0	0	0	0
11	0	0	0	0	0	0	0	0	0	0
12	0	0	0	1	0	0	0	0	0	0
13	0	0	0	0	0	0	0	0	1	0
14	0	0	0	0	0	0	0	1	0	0
15	0	0	0	0	1	0	0	1	0	0
16	0	0	0	0	0	0	0	0	0	0
17	0	0	0	0	0	0	0	0	0	0
18	0	0	0	0	0	0	0	1	0	0
19	0	0	0	0	0	0	0	0	1	0
20	0	1	0	0	0	0	0	1	0	0
21	0	0	0	0	1	0	0	0	0	1


- Marketing - Ads Selection
 - Random Selection

```
1 # Random Selection
2
3 # Importing the libraries
4 import numpy as np
5 import matplotlib.pyplot as plt
6 import pandas as pd
7
8 # Importing the dataset
9 dataset = pd.read_csv('Ads_CTR_Optimisation.csv')
10
11 # Implementing Random Selection
12 import random
13 N = 10000
14 d = 10
15 ads_selected = []
16 total_reward = 0
17 for n in range(0, N):
18     ad = random.randrange(d)
19     ads_selected.append(ad)
20     reward = dataset.values[n, ad]
21     total_reward = total_reward + reward
22
23 # Visualising the results - Histogram
24 plt.hist(ads_selected)
25 plt.title('Histogram of ads selections')
26 plt.xlabel('Ads')
27 plt.ylabel('Number of times each ad was selected')
28 plt.show()
```

Marketing - Ads Selection

○ Random Selection

j ▲	Type	Size	Value
0	int	1	4
1	int	1	6
2	int	1	1
3	int	1	0
4	int	1	4
5	int	1	5
6	int	1	0
7	int	1	8
8	int	1	5
9	int	1	3
10	int	1	8
11	int	1	5
12	int	1	8
13	int	1	4
14	int	1	0




total_reward

int64


1

1193

Upper Confidence Bound vs Thomson Sampling




UCB



- Deterministic
- Requires update at every round

Thompson Sampling



- Probabilistic
- Can accommodate delayed feedback
- Better empirical evidence

■ Upper Confidence Bound Algorithm

Step 1. At each round n , we consider two numbers for each ad i :

- $N_i(n)$ - the number of times the ad i was selected up to round n ,
- $R_i(n)$ - the sum of rewards of the ad i up to round n .

Step 2. From these two numbers we compute:

- the average reward of ad i up to round n

$$\bar{r}_i(n) = \frac{R_i(n)}{N_i(n)}$$

- the confidence interval $[\bar{r}_i(n) - \Delta_i(n), \bar{r}_i(n) + \Delta_i(n)]$ at round n with

$$\Delta_i(n) = \sqrt{\frac{3 \log(n)}{2 N_i(n)}}$$

Step 3. We select the ad i that has the maximum UCB $\bar{r}_i(n) + \Delta_i(n)$.


- Marketing - Ads Selection
 - Upper Confidence Bound

```
1 # Upper Confidence Bound
2
3 # Importing the libraries
4 import numpy as np
5 import matplotlib.pyplot as plt
6 import pandas as pd
7
8 # Importing the dataset
9 dataset = pd.read_csv('Ads_CTR_Optimisation.csv')
10
11 # Implementing UCB
12 import math
13 N = 10000
14 d = 10
15 ads_selected = []
16 numbers_of_selections = [0] * d
17 sums_of_rewards = [0] * d
18 total_reward = 0
19 for n in range(0, N):
20     ad = 0
21     max_upper_bound = 0
22     for i in range(0, d):
23         if (numbers_of_selections[i] > 0):
24             average_reward = sums_of_rewards[i] / numbers_of_selections[i]
25             delta_i = math.sqrt(3/2 * math.log(n + 1) / numbers_of_selections[i])
26             upper_bound = average_reward + delta_i
27         else:
28             upper_bound = 1e400
29         if upper_bound > max_upper_bound:
30             max_upper_bound = upper_bound
31             ad = i
32     ads_selected.append(ad)
33     numbers_of_selections[ad] = numbers_of_selections[ad] + 1
34     reward = dataset.values[n, ad]
35     sums_of_rewards[ad] = sums_of_rewards[ad] + reward
36     total_reward = total_reward + reward
37
```

Marketing - Ads Selection

- Upper Confidence Bound

i	Type	Size	Value
9985	int	1	4
9986	int	1	4
9987	int	1	4
9988	int	1	4
9989	int	1	4
9990	int	1	4
9991	int	1	4
9992	int	1	4
9993	int	1	4
9994	int	1	4
9995	int	1	4
9996	int	1	4
9997	int	1	4
9998	int	1	4
9999	int	1	4



total_reward	int64	1	2178
upper_bound	float64	1	0.31017236647899182

■ Thompson Sampling Algorithm

Step 1. At each round n , we consider two numbers for each ad i :

- $N_i^1(n)$ - the number of times the ad i got reward 1 up to round n ,
- $N_i^0(n)$ - the number of times the ad i got reward 0 up to round n .

Step 2. For each ad i , we take a random draw from the distribution below:

$$\theta_i(n) = \beta(N_i^1(n) + 1, N_i^0(n) + 1)$$

Step 3. We select the ad that has the highest $\theta_i(n)$.

■ Thompson Sampling Algorithm

- Ad i gets rewards \mathbf{y} from Bernoulli distribution $p(\mathbf{y}|\theta_i) \sim \mathcal{B}(\theta_i)$.
- θ_i is unknown but we set its uncertainty by assuming it has a uniform distribution $p(\theta_i) \sim \mathcal{U}([0, 1])$, which is the prior distribution.
- Bayes Rule: we approach θ_i by the posterior distribution

$$\underbrace{p(\theta_i|\mathbf{y})}_{\text{posterior distribution}} = \frac{p(\mathbf{y}|\theta_i)p(\theta_i)}{\int p(\mathbf{y}|\theta_i)p(\theta_i)d\theta_i} \propto \underbrace{p(\mathbf{y}|\theta_i)}_{\text{likelihood function}} \times \underbrace{p(\theta_i)}_{\text{prior distribution}}$$

- We get $p(\theta_i|\mathbf{y}) \sim \beta(\text{number of successes} + 1, \text{number of failures} + 1)$
- At each round n we take a random draw $\theta_i(n)$ from this posterior distribution $p(\theta_i|\mathbf{y})$, for each ad i .
- At each round n we select the ad i that has the highest $\theta_i(n)$.

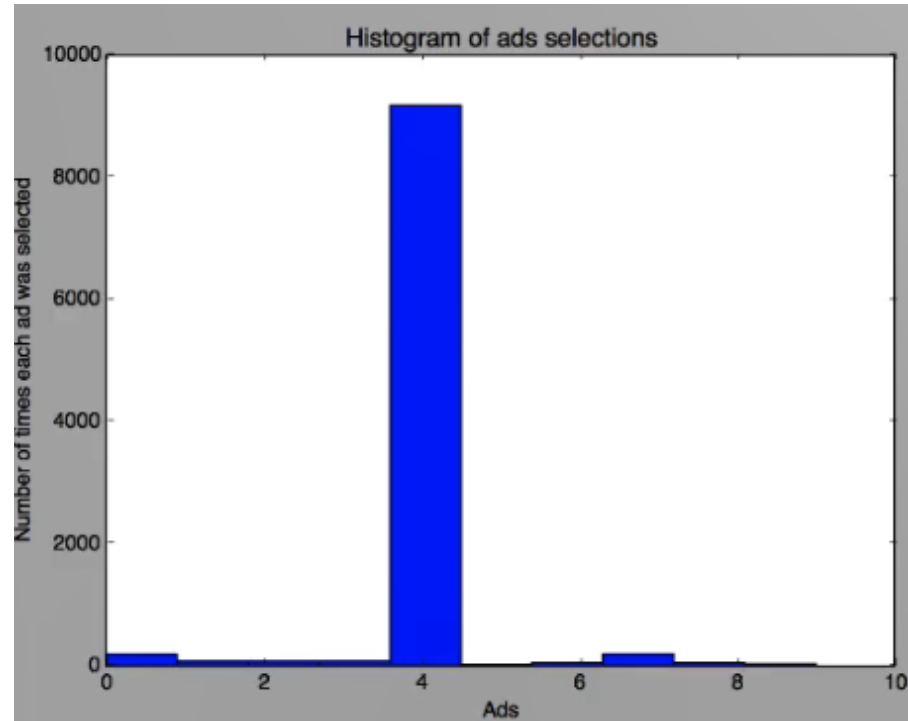
- Marketing - Ads Selection
 - Thompson Sampling

```
1# Thompson Sampling
2
3# Importing the libraries
4import numpy as np
5import matplotlib.pyplot as plt
6import pandas as pd
7
8# Importing the dataset
9dataset = pd.read_csv('Ads_CTR_Optimisation.csv')
10
11# Implementing Thompson Sampling
12import random
13N = 10000
14d = 10
15ads_selected = []
16numbers_of_rewards_1 = [0] * d
17numbers_of_rewards_0 = [0] * d
18total_reward = 0
19for n in range(0, N):
20    ad = 0
21    max_random = 0
22    for i in range(0, d):
23        random_beta = random.betavariate(numbers_of_rewards_1[i] + 1, numbers_of_rewards_0[i] + 1)
24        if random_beta > max_random:
25            max_random = random_beta
26            ad = i
27    ads_selected.append(ad)
28    reward = dataset.values[n, ad]
29    if reward == 1:
30        numbers_of_rewards_1[ad] = numbers_of_rewards_1[ad] + 1
31    else:
32        numbers_of_rewards_0[ad] = numbers_of_rewards_0[ad] + 1
33    total_reward = total_reward + reward
34
```

■ Marketing - Ads Selection

○ Thompson Sampling

i	Type	Size	Value
9985	int	1	4
9986	int	1	4
9987	int	1	4
9988	int	1	4
9989	int	1	4
9990	int	1	4
9991	int	1	4
9992	int	1	4
9993	int	1	4
9994	int	1	4
9995	int	1	4
9996	int	1	4
9997	int	1	4
9998	int	1	4
9999	int	1	4



total_reward	int64	1	2613
--------------	-------	---	------



Universidade do Minho
Escola de Engenharia
Departamento de Informática

Mestrado Integrado em Engenharia Informática

Computação Natural

2018/2019

Paulo Novais, Cesar Analide, Filipe Gonçalves